



Linux Containers From Scratch

Joshua Hoffman

Velocity Europe 2014

SETUP

INSTALL PACKAGES

Recommended mirror:

<http://ftp.es.debian.org>

Install packages:

- vim
- screen
- lftp
- busybox-static
- systemd
- yum
- qemu-utils
- aufs-tools
- pbzip2
- htop

SETUP

CONFIGURE SYSTEMD

1. Edit /etc/default/grub

change the line:

```
GRUB_CMDLINE_LINUX=""
```

to:

```
GRUB_CMDLINE_LINUX="init=/bin/systemd"
```

2. Run the grub updater:

```
update-grub2
```

3. Reboot

THE CLOUD LINUX CONTAINERS

THE CLOUD LINUX CONTAINERS FREE LUNCH

DO NOT EXIST

**IDEAS
NOT
THINGS**

PORTABILITY

ISOLATION

VIRTUAL ~~MACHINE~~ ENVIRONMENT

A logically isolated virtual environment.

A Linux Container

**FUNDAMENTALLY
DIFFERENT THAN
VIRTUAL MACHINES**

TRANSPARENT

Running in a Virtual Machine

as viewed from the host os

```
# ps x
```

PID	TTY	STAT	TIME	COMMAND
689	?	R	1:06	qemu-kvm

Running in a Linux Container

as viewed from the host os

```
# ps x
```

```
PID  TTY  STAT TIME  COMMAND
```

```
5347 ?    R      2:22 unicorn_rails master -D -c kiffen.rb
```

NAMESPACES

NAMESPACES: NETWORK

NETWORK NAMESPACE

as viewed from iproute2

```
$ ip a
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 16436 qdisc noqueue state  
UNKNOWN
```

```
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

```
    inet 127.0.0.1/8 scope host lo
```

```
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc  
pfifo_fast master br0 state UP qlen 1000
```

```
    link/ether 00:01:2e:3b:be:14 brd ff:ff:ff:ff:ff:ff
```

```
    inet 10.21.0.22/24 brd 10.21.0.255 scope global br0
```

```
    inet6 fe80::201:2eff:fe3b:be14/64 scope link
```

```
        valid_lft forever preferred_lft forever
```

NAMESPACES: MOUNT

MOUNT NAMESPACE

as viewed from ls

```
$ ls /
```

```
bin    etc      lib      media    proc     sbin     sys      var
boot   home     lib64    mnt      root     selinux  tmp
dev    lost+found  opt      run      srv      usr
```

NAMESPACES:

PID

PID NAMESPACE

as viewed from ps

```
# ps x
```

```
PID  TTY  STAT TIME  COMMAND
```

```
5347 ?    R      2:22 unicorn_rails master -D -c kiffen.rb
```

CGROUPS

CGROUPS

as viewed from ls

```
# ls -F /sys/fs/cgroup/  
blkio/  cpu@  cpuacct@cpu,cpuacct/  cpuset/  devices/  
freezer/  net_cls/  perf_event/  systemd/
```

```
# ls -F /sys/fs/cgroup/cpuset  
cpuset.mem_exclusive  cgroup.procs  
cpuset.memory_migrate cpuset.mems  
cpuset.cpu_exclusive  tasks  cpuset.cpus  
(...output truncated...)
```


DEMO:

exploring containers

with busybox

Minimal Busybox Container

```
# mkdir -p {minimal,minimal/usr}/{bin,sbin,etc}
```

```
# for x in $(busybox --list-full); do  
> ln -s /bin/sh minimal/$x; done
```

```
# cp -f /bin/busybox minimal/bin/sh
```

```
# touch minimal/etc/os-release
```

Running The Container

Private mount namespace:

```
# chroot minimal /bin/sh
```

Private mount and pid namespace

```
# systemd-nspawn -Dminimal /bin/sh
```

Private mount, pid, and network namespace

```
# systemd-nspawn --private-network -Dminimal /bin/sh
```

DEMO:

building a container

image with cpio

Build A Container Image With cpio

```
# find minimal -print | cpio -o |  
> pbzip2 -c > minimal.cpio.bz2
```

```
# ls -lh minimal.cpio.bz2  
-rw-r--r--  1 root root 852K Nov 18 12:48 minimal.cpio.bz2
```

DEMO:

limiting cpu access

with cgroups

Limiting CPU Access With cgroups

```
# dd if=/dev/urandom of=datafile bs=1M count=100
```

```
# time pbzip2 -k -9 datafile
```

```
# mkdir /sys/fs/cgroup/cpuset/my_cpuset
```

```
# echo 0 > /sys/fs/cgroup/cpuset/my_cpuset/cpuset.cpus
```

```
# echo 0 > /sys/fs/cgroup/cpuset/my_cpuset/cpuset.mems
```

```
# echo $$ > /sys/fs/cgroup/cpuset/my_cpuset/tasks
```

```
# time pbzip2 -k -9 datafile
```

DEMO:

connect a container

to the network

Connect The Network With iproute2

```
# ip netns add minimal
# ip link add eth1 type veth peer name veth1
# ip link set eth1 netns minimal
# ip a add 10.0.0.1/24 dev veth1
# ip l set veth1 up
# ip netns exec minimal chroot minimal /bin/sh
```

(in the container)

```
# ip a add 10.0.0.2/24 dev eth1
# ip l set eth1 up
```

DEMO:

installing a service stack with yum

SETUP

CONFIGURE YUM

Create a file called yum.conf with the following contents:

```
[main]
cachedir=/var/cache/yum
keepcache=1
debuglevel=2
logfile=/var/log/yum.log
exactarch=1
obsoletes=1

[base]
name=CentOS-7 - Base
#mirrorlist=http://mirrorlist.centos.org/?release=7&arch=x86_64&repo=os
baseurl=http://192.168.56.1/centos/
gpgcheck=0
enabled=1
```

Install A Service Stack With yum

```
# mkdir -p /lcfs/ftp_stack
# yum -c yum.conf --installroot=/lcfs/ftp_stack \
> install vsftpd

# ip netns exec minimal chroot /lcfs/ftp_stack /bin/bash

(in the container)
# /sbin/vsftpd
```

DEMO:

splitting a container

image into layers with

aufs

Container Layers With aufs

```
# mkdir -p /lcfs/base_stack
# yum -c yum.conf \
> --installroot=/lcfs/base_stack install basesystem
# cp yum.conf /lcfs/base_stack/etc/
# rm /lcfs/base_stack/etc/yum.repos.d/*repo

# mkdir /lcfs/{app_stack,tmp_stack}
# mount -t aufs -obr=/lcfs/app_stack:/lcfs/base_stack none \
> /lcfs/tmp_stack

# yum --installroot=/lcfs/tmp_stack install vsftpd
```

DEMO:

install a full os with

yum

Install A Full OS With yum

```
# mkdir -p /lcfs/centos-rootfs  
# yum -c yum.conf --installroot=/lcfs/centos-rootfs \  
> groupinstall core
```

```
# chroot /lcfs/centos-rootfs  
# passwd (set a new password)
```

```
# vi /etc/pam.d/session (comment these out lines)  
session      required      pam_selinux.so close  
session      required      pam_loginuid.so  
session      required      pam_selinux.so open
```


Run A Full OS Container

```
# systemd-nspawn --private-network -D/lcfs/centos-rootfs
```

